

Intro, C refresher

CSE 333 Winter 2021

Instructor: John Zahorjan

Welcome – please set up
your Zoom session. We'll
start the actual class meeting
at 11:30 am pdt

Teaching Assistants:

Matthew Arnold	Nonthakit Chaiwong	Jacob Cohen
Elizabeth Haker	Henry Hung	Chase Lee
Leo Liao	Tim Mandzyuk	Benjamin Shmidt
Guramrit Singh		

Lecture Outline

- ❖ **Course Introduction**
- ❖ Course Policies
 - <https://cs.washington.edu/333> / Syllabus tab
- ❖ Course Map
- ❖ C Intro

To get started...

- ❖ It's all remote, all the time this quarter
- ❖ Core infrastructure is same as always
(Gradescope, Gitlab, web, discussion board)
- ❖ Lectures, sections, office hours – Zoom
- ❖ Stay healthy in every way

Lectures

- ❖ Classes will be mostly lectures – more interaction in sections
 - Varied experiences so far. Let us know where we could do better!
- ❖ Conventions (from page on our web site)
 - Lecture will be recorded and available to class members (only)
 - Recordings are kept only a short while past the end of the quarter
 - If you have a question, (a) thank you!, and (b) type it in the chat window or speak up
 - We intend to post lecture slides in advance

Online Sections

❖ Sections: more Zoom

- Not normally recorded so we can have open discussions and group work without people being too self-conscious
- We're going to try to produce videos for things that would normally be done as demos or presentations; details TBA
 - Those will be available online
- Slides and any sample code, worksheets, etc. posted as always

Online Everything Else

- ❖ Office hours: also Zoom; combination of group gatherings, breakouts, waiting rooms, sign-up sheets to organize – all as needed
 - Not recorded or archived
 - Once gitlab repos are set up, if your question concerns your code (exercises, projects), please push latest code to the repo before meeting with TA to save some time
- ❖ You will be bombarded with email as we add these things to Canvas/Zoom. Feel free to file away for future reference. 😊

Stay in Touch – Speak up

- ❖ This is a strange world we're in and there's a lot of stress for many people
- ❖ Please speak up if things aren't (or are!) going well
 - We can often help if we know about things, so stay in touch with TAs, instructor, advising, friends and peers, others
- ❖ We're all in this together but not all in the same way, so please show understanding and help us understand

Introductions: Course Staff

- ❖ John Zahorjan (instructor)
 - Long-time CSE faculty member and CSE 333 veteran
- ❖ TAs:
 - Matthew Arnold, Nonthakit Chaiwong, Jacob Cohen, Elizabeth Haker, Henry Hung, Chase Lee, Leo Liao, Tim Mandzyuk, Benjamin Shmidt, Guramrit Singh
 - Available in section, office hours, and discussion group
 - An invaluable source of information and help
- ❖ Get to know us
 - We are here to help you succeed!

Introductions: Students

- ❖ ~135 students this quarter
 - There are no overload forms or waiting lists for CSE courses
- ❖ Time zones?
- ❖ Expected background
 - **Prereq:** CSE 351 – a little bit of hardware architecture (registers, memory, instructions), a little bit of C (notion of types, pointers, procedure call convention, compile/link build)
 - CSE 391 or Linux skills needed for CSE 351 helpful, but not having that isn't a show stopper

Lecture Outline

- ❖ Course Introduction
- ❖ **Course Policies**
 - <https://courses.cs.washington.edu/courses/cse333/21wi/syllabus/>
 - Summary here, but you **must** read the full details online
- ❖ Course Map
- ❖ C Intro

Communication

- ❖ **Website:** <http://cs.uw.edu/333>
 - Schedule, policies, materials, assignments, etc.
- ❖ **Discussion:** Ed group linked to course home page
 - Must log in using your **@uw.edu** Google identity (not cse)
 - Ask and answer questions – staff will monitor and contribute
- ❖ **Staff mailing list:** cse333-staff@cs for things not appropriate for Ed group
 - Mail is sent to instructor and all TAs
- ❖ **Course mailing list:** for announcements from staff
 - Registered students automatically subscribed with your @uw email
- ❖ **Office Hours:** spread throughout the week
 - Schedule posted shortly and will start right away
 - Can also e-mail to staff list to make individual appointments
 - Will try to consider time zones when scheduling

Course Components

- ❖ Lectures (28)
 - Introduce the concepts
- ❖ Sections (10)
 - Applied concepts, important tools and skills for assignments, clarification of lectures, exam review and preparation
- ❖ Programming Exercises (~17)
 - Roughly one per lecture, due the morning before the next lecture
 - Coarse-grained grading (0, 1, 2, or 3)
- ❖ Programming Projects (0+4)
 - Warm-up, then 4 “homeworks” that build on each other
- ❖ No traditional exams, but hoping to do ~4 “recap/review” assignments for things traditionally covered on exams

Grading (tentative)

- ❖ **Exercises:** ~35%
 - Submitted via GradeScope (account info mailed this morning)
 - Graded on correctness and beauty by TAs
- ❖ **Projects:** ~45% total
 - Submitted via GitLab; must tag commit that you want graded
 - Binaries provided if you didn't get previous part working
- ❖ **Note the lack of quiz and exam points**

Deadlines and Student Conduct

❖ Official late policy

- Exercises: no late submissions accepted, due 10 am
- Projects: 4 late days for entire quarter, max 2 per project
- Need to get things done on time – difficult to catch up!
 - But given remote world, we'll work with you if things come up

❖ Academic Integrity (read the full policy on the web)

- I trust you implicitly and will follow up if that trust is violated
- In short: don't attempt to gain credit for something you didn't do and don't help others do so either
- This does **not** mean suffer in silence – learn from the course staff and peers, talk, share ideas; *but* don't share or copy work that is supposed to be yours

And off we go...

- ❖ Goal is to figure out setup and computing infrastructure right away so we don't put that off and then have a crunch later in the quarter
- ❖ So:
 - First exercise out today, due Friday morning **10 am**
 - **First homework out Wednesday, due next Monday**
 - Warmup/logistics in sections Thursday
 - HW0 (the warmup project) published this afternoon and gitlab repo's created.

Okay, but...

- ❖ CSE 333 has a tightly integrated set of projects, exercises, and lectures
- ❖ Okay, but...
- ❖ I think things are hard enough, and I'm naturally inclined to think that having to worry about the details makes it harder to reflect on anything larger
- ❖ So... the mantra for this quarter is “flexibility”
 - Grading
 - Due dates
 - Programming style rules

Something old

- ❖ Same homeworks...
 - Students who have done them are generally enthusiastic about them
- ❖ Mainly most same exercises
 - This is just fear – I don't know at this point how tightly integrated the exercises, lectures, and homeworks are
 - I endorse the idea of frequent, small exercises, though
- ❖ This quarters slides will be derived from “the standard slides for CSE 333”

Something new

- ❖ Today's class isn't exactly what usually happens
- ❖ My intention is for this class to be “gentle”
 - Carrots
- ❖ My intention is to help you learn more than you would if you were left to learn this material on your own
- ❖ I've monkeyed with hw0 / ex00
- ❖ I've made some policy changes
 - Meet A TA sessions
 - Homeworks
 - Homeworks 0 – 2 (in C) done individually
 - Homework 3 (C++) done in pairs
 - Home 4 (C++) done in different pairs
 - We suggest “pair programming”

Meet A TA Session

❖ Brief...

❖ Your part starts with:

""Hi, _____, my name is _____. I'm a fresh/soph/junior/senior majoring in _____. My favorite CSE course so far was _____. The thing I liked about it was _____. My favorite course ever, anywhere, was _____. The thing I liked about it was _____. I'm [really/kind of] looking forward to _____ in CSE 333 this quarter. I'm worried about _____ in CSE 333 this quarter. [*Optional: Something I'd like you to know about me is _____.*] I have this question: _____."

Team Programming

- ❖ Traditionally, all work done individually
 - All the infrastructure seems to be affected by that
- ❖ We're going where no CSE 333 has gone before
 - Who knows what can go wrong!
- ❖ *Aside: I've changed the version of C/C++ we're using from c11/c++11 to c17/c++17, but it turns out there's a course tool that might depend on c11/c++11...*
- ❖ This quarter we'll do some projects individually, some in pairs
- ❖ Why?
 - Programming alone is an anachronism
 - Every classmate knows something you don't, and vice versa
 - C is pretty tame – questions tend to be simpler, and web searches more effective
 - C++ is a monster – I think it's very likely having someone to talk over details of some implementation issue will be good for almost everyone

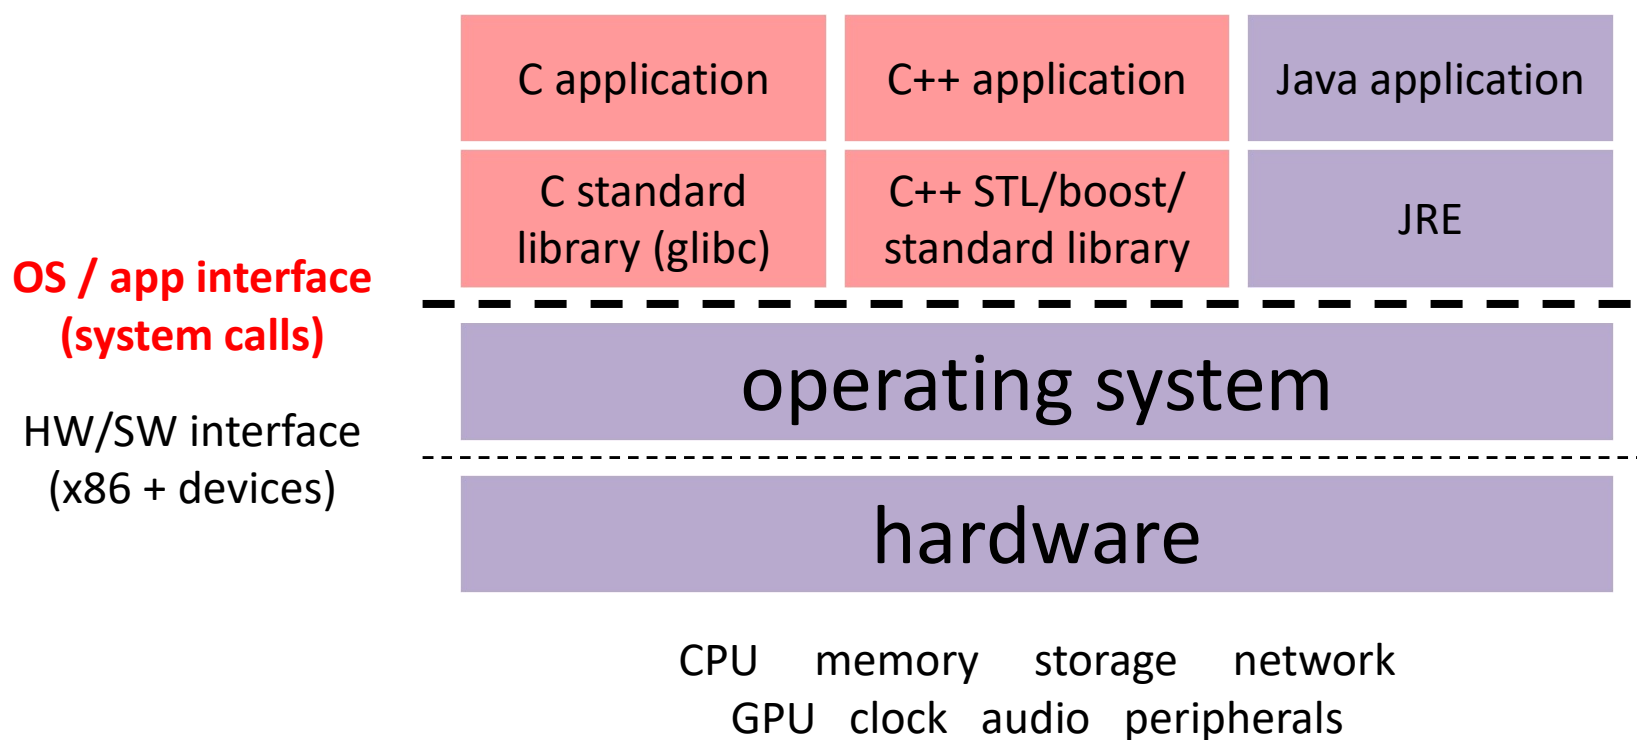
Changing Teams Mid-Stream

- ❖ For all we know, this could be a great idea...
- ❖ But homeworks are cumulative...
- ❖ Your goal:
 - Be a teammate everyone would want to work with
 - Learn something
- ❖ Not your goal:
 - Producing the finest project cse333 has ever seen
- ❖ If you're capable of doing a better implementation on your own than your team can produce, then set your goal on figuring out how to make it so you aren't
- ❖ Grades...
 - Effort counts

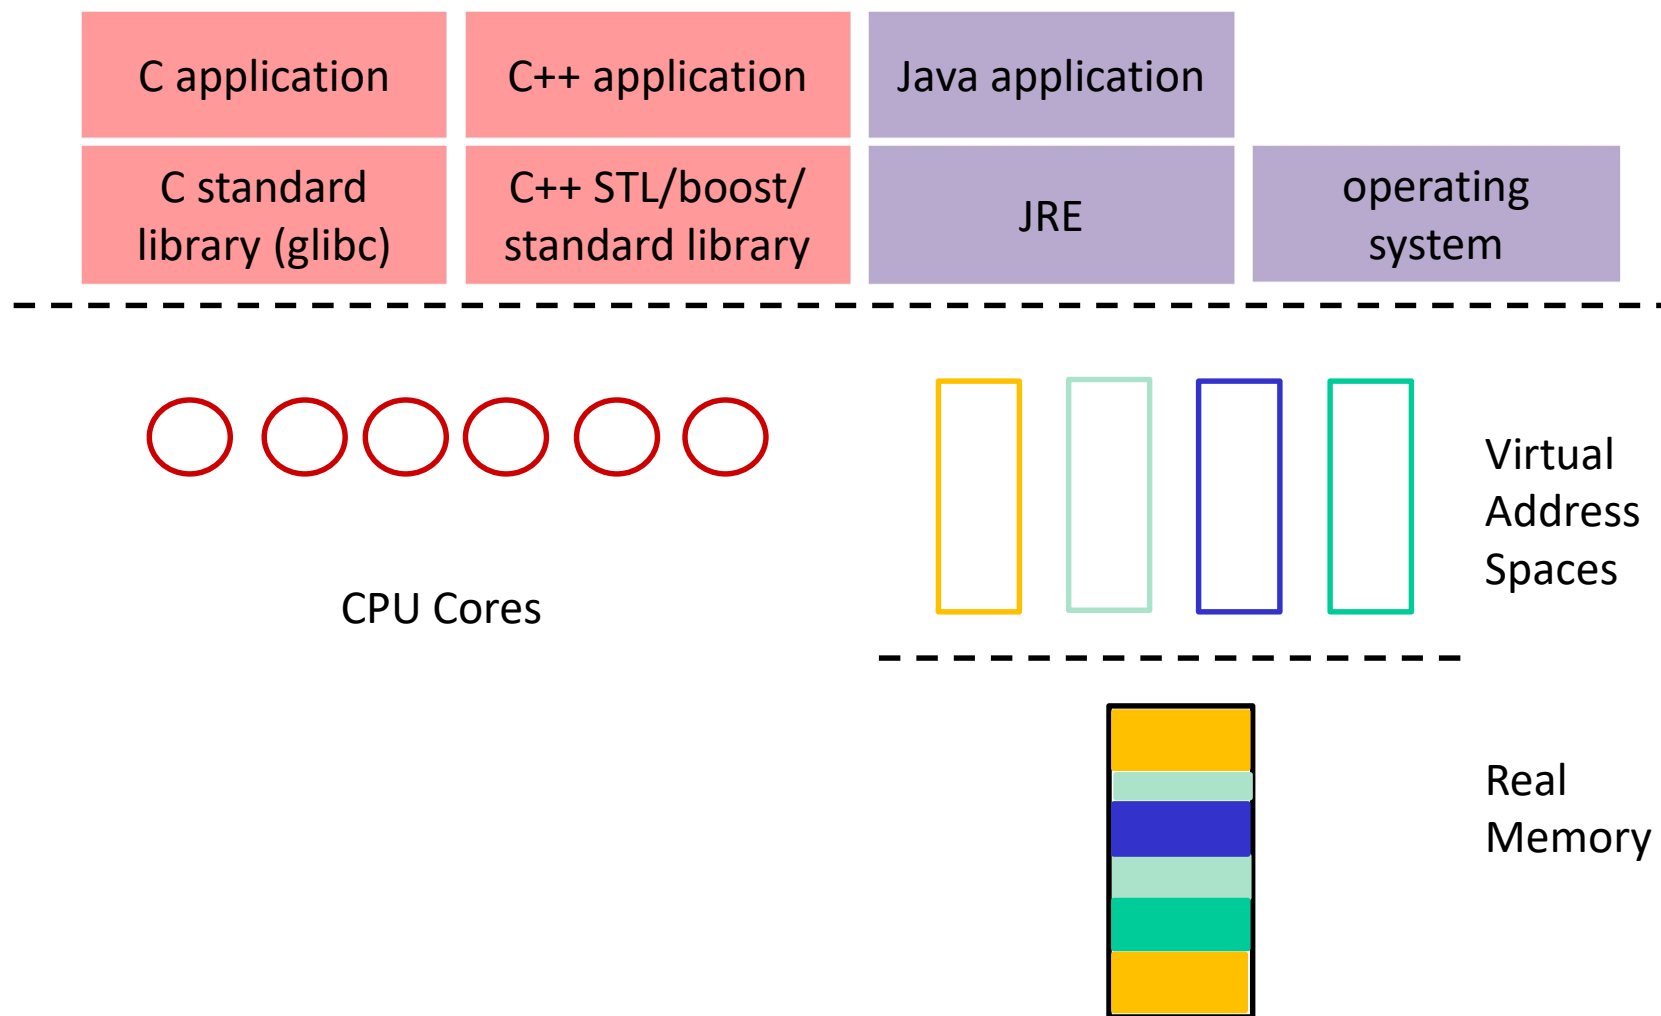
Lecture Outline

- ❖ Course Introduction
- ❖ Course Policies
 - <https://cs.washington.edu/333> / Syllabus tab
- ❖ **Course Map**
- ❖ C Intro

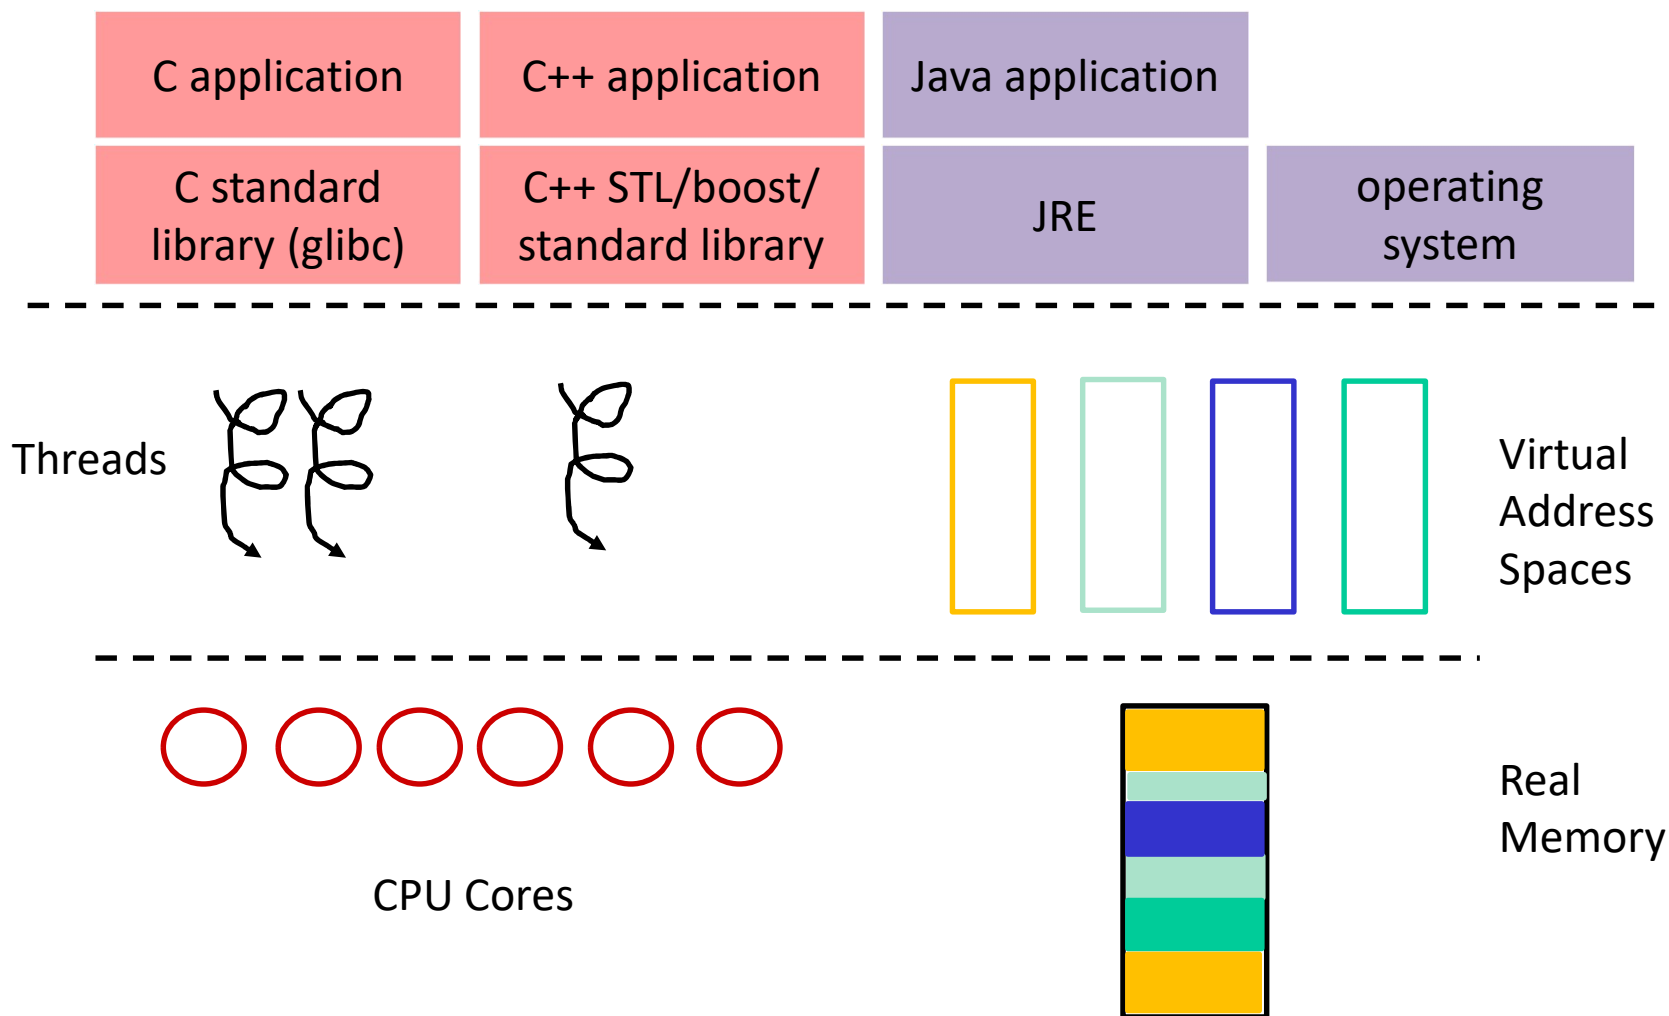
Course Map: 100,000 foot view



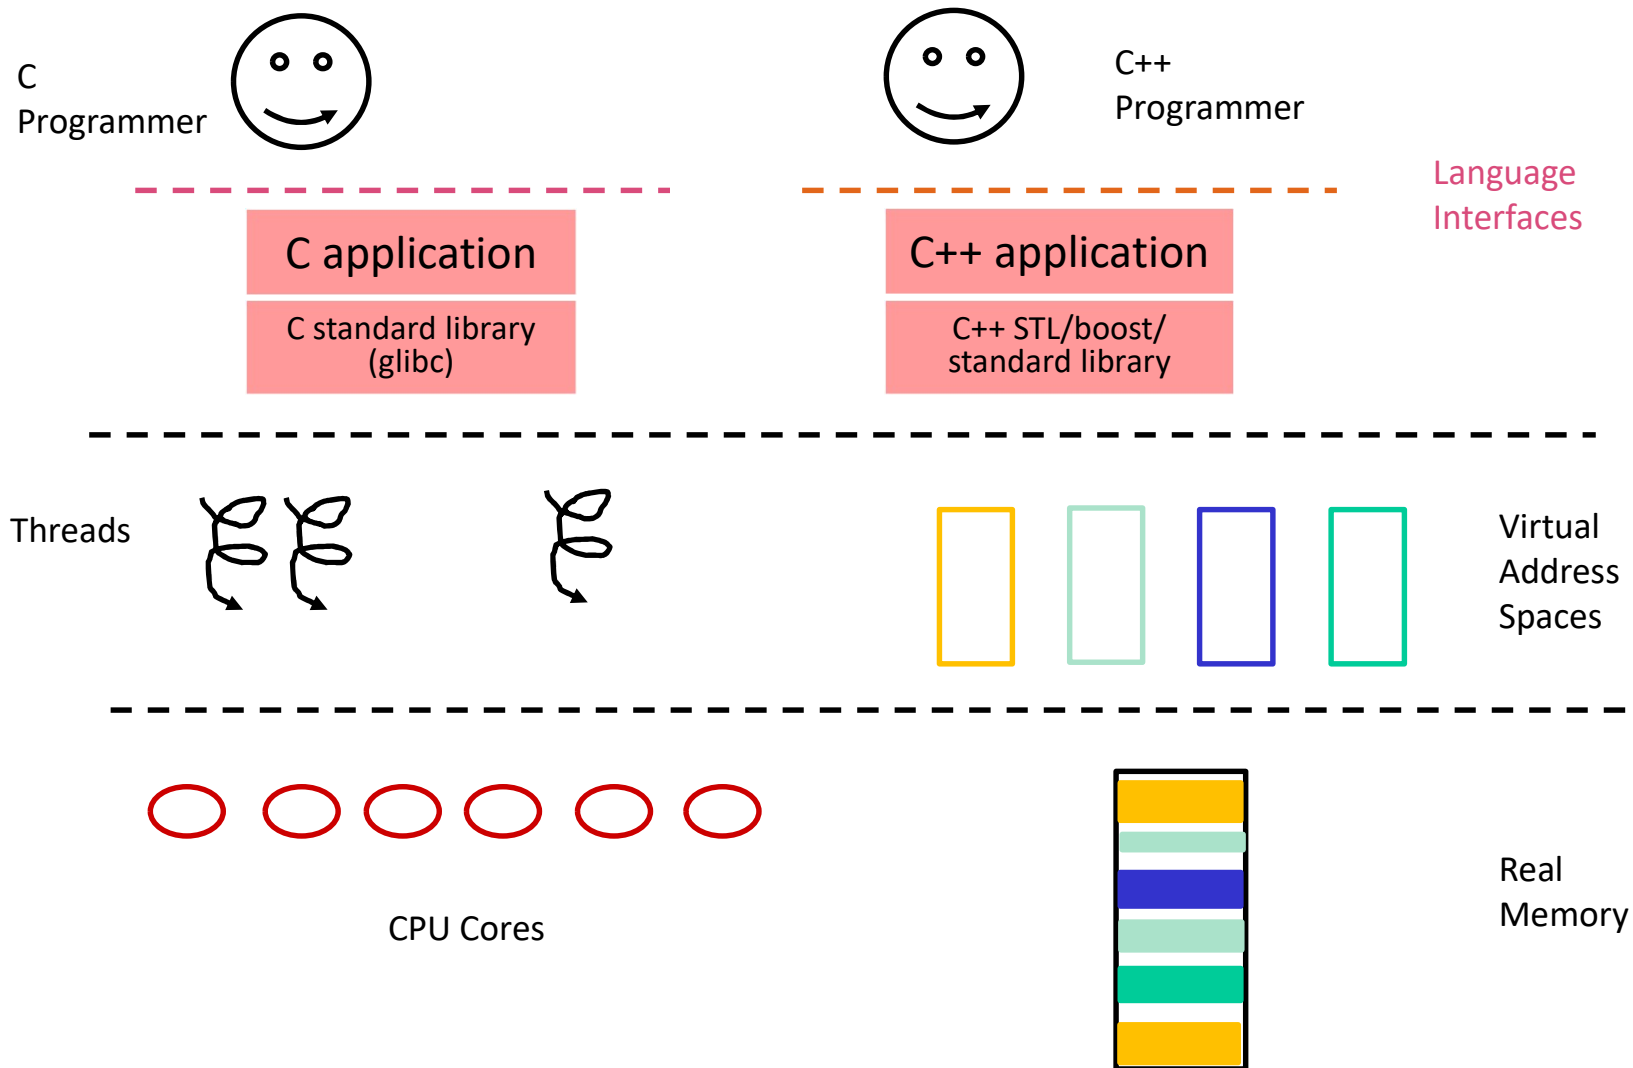
Course Map Picture Revisited I



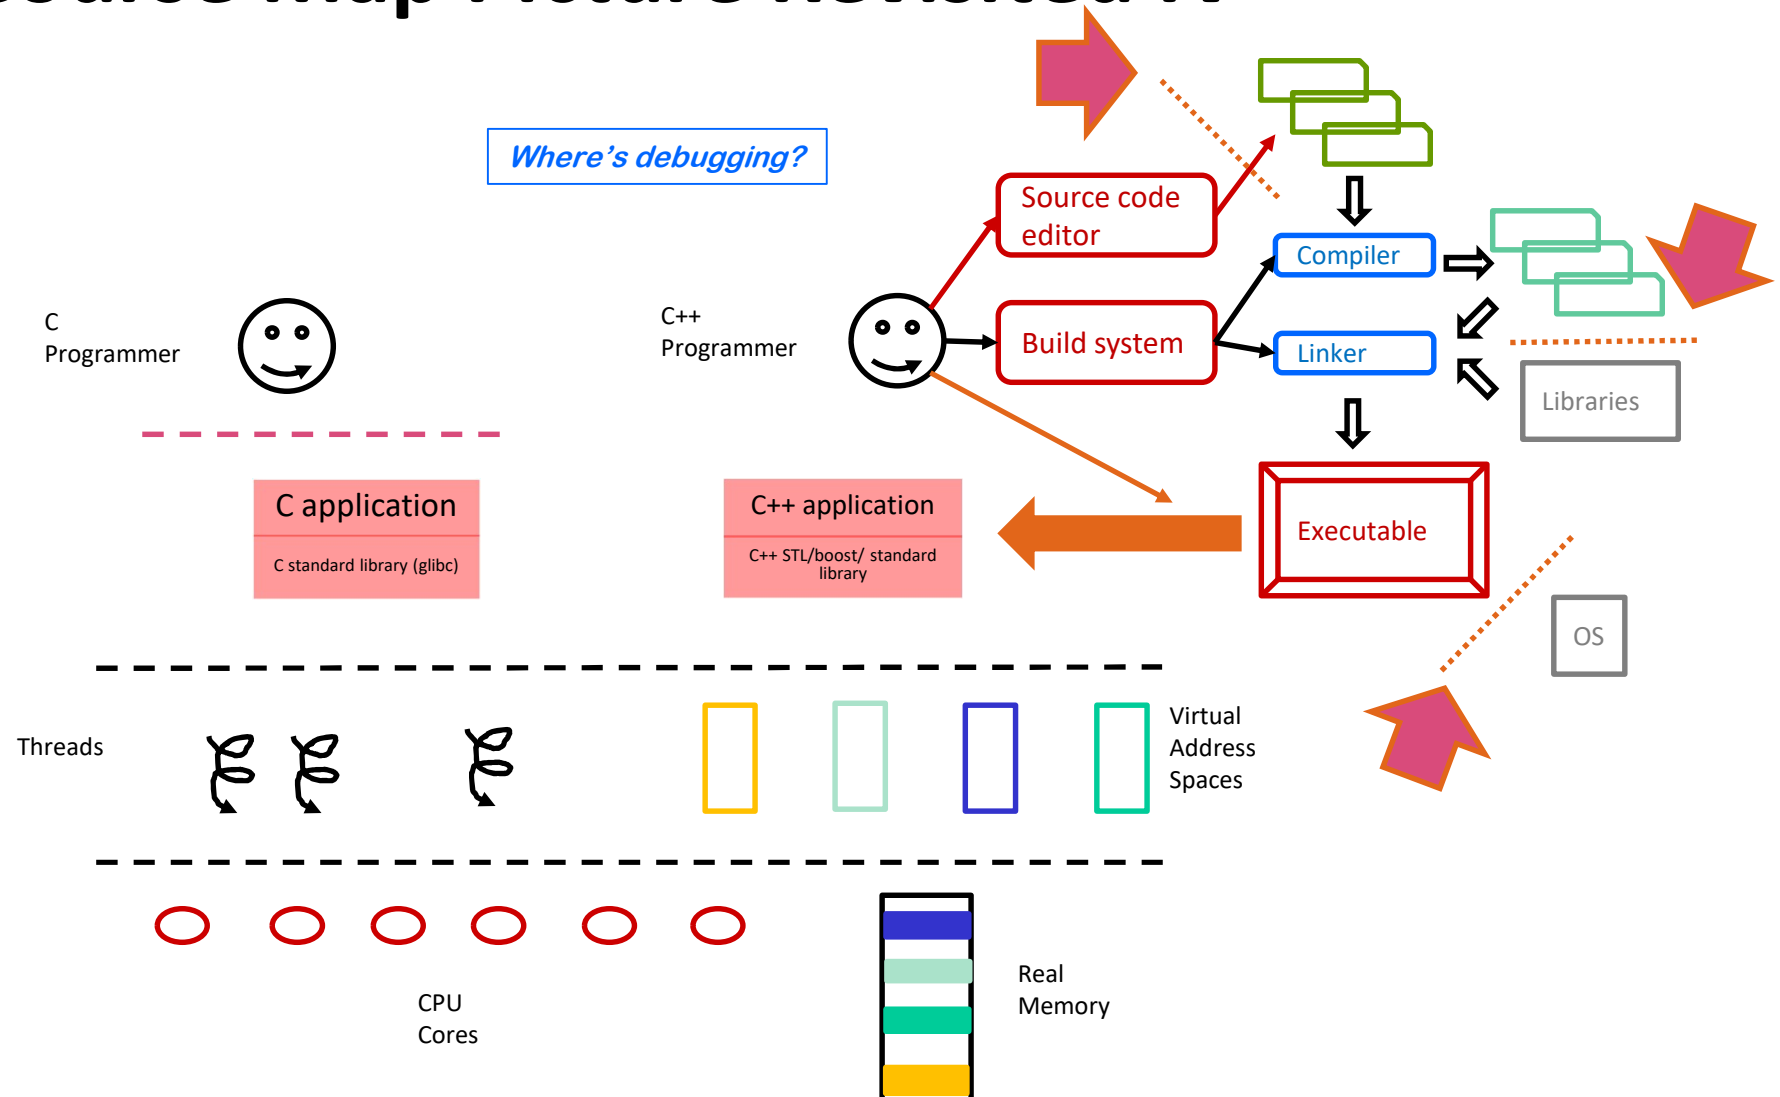
Course Map Picture Revisited II



Course Map Picture Revisited III



Course Map Picture Revisited IV



CSE 333 21wi Tools

Static

- ❖ OS
 - CentOS 8.2 (attu.cs.washington.edu)
- ❖ Compilers / libraries
 - gcc 9.2.1 / g++ 9.2.1
- ❖ Build
 - GNU make 4.2.1
- ❖ Language versions
 - C 2017 / C++ 2017
- ❖ Editors
 - vscode
 - vim
 - emacs
- ❖ Source Control
 - git / gitlab

Dynamic

- ❖ Debugger
 - gdb
- ❖ Unit Test
 - Gtest
- ❖ Other Tests
 - valgrind

Systems Programming

❖ “Systems”

- Not an application
- Or perhaps an application that facilitates building other applications

❖ “Programming”

- In CSE 333, in C/C++ -- systems languages
- In CSE 333, “programming” includes the tools/procedures to go from source code to debugged execution
 - The tools are often language specific
 - Many of the concepts are not

Lecture Outline

- ❖ Course Introduction
- ❖ Course Policies
 - <https://courses.cs.washington.edu/courses/cse333/18sp/syllabus/>
- ❖ Course Map
- ❖ **C Intro**
 - **Workflow, Variables, Functions**


C

❖ Created in 1972 by Dennis Ritchie

- Designed for creating system software
- Portable across machine architectures
- Most recently updated in 1999 (C99) and 2011 (C11) and 2017 (C17)
 - Currently working on C 2x

❖ Characteristics

- “Low-level” language that allows us to exploit underlying features of the architecture – **but easy to fail spectacularly (!)**
- Procedural (not object-oriented)
- Typed but unsafe (possible to bypass the type system)
- Small, basic library compared to Java, C++, most others....



Understanding C

- ❖ Assembler gives access to “everything” on the processor
- ❖ Assembler programs mainly involve:
 - Deciding how to use memory
 - How much memory is required to store a value
 - What values live where
 - How does program address those values at run time
 - Example: a logical array of values stored in consecutive memory locations
 - Deciding what values to cache in registers, when
 - Writing instructions that operate on values
 - Example: expression evaluation
 - Implementing control flow
 - Examples: loops; procedure call and return
- ❖ Don't think of assembler/machine code as how C programs are realized, think of C as a more convenient way to express an assembler program

What Does the C Compiler Do?

- ❖ Certain tasks the assembler programmer would have to do can be automated
 - Names, not addresses
 - myInt, not 0(eax)
 - Memory size
 - How much space is needed for an int? a float? a char?
 - Memory layout
 - If a procedure has three integer local variables, what are their offsets from the frame pointer?
 - Control flow instructions
 - for/while loops
 - Procedure call/return
 - Expression evaluation
 - Example: $7+x*(y-\text{sub}(z))$

Compiling C vs. Assembling Assembly

- ❖ It seems natural to think of programs as simply a specification of an execution
 - That is, we make little mental distinction between the static code file and the execution of the program
- ❖ In assembler, nothing very interesting happens at assembly time
 - `mov 0x0(%rip),%rax => 48 8b 05 00 00 00 00`
- ❖ In contrast, what does the C compiler have to do with these?
 - `z = "One" + (long int)"Two";`
`y = x * Five - process(val);`
`y /= x unless x==0; // Note: This isn't C... But why not?`

C++ vs. C

C17 vs. C99

- ❖ In a sense, the “action” in these languages has to do with things that happen at compile time
 - For the most part, what can be done at run time hasn’t changed much
 - The language helps you write what you mean succinctly
 - It helps you write robust, correct code
 - It helps keep you from writing incorrect code
 - It helps you write performant code
 - [You help it generate performant code]
- ❖ C++ example:
 - `auto val = CreateValue(y);` // Declaration of val. Compiler
// determines correct type

Generic C Program Layout

```
#include <system_files>
#include "local_files"

#define macro_name macro_expr

/* declare functions */
/* declare external variables & structs */

int main(int argc, char* argv[]) {
    /* the innards */
}

/* define other functions */
```

C Syntax: `main`

- ❖ To get command-line arguments in `main`, use:

```
int main(int argc, char* argv[])
```

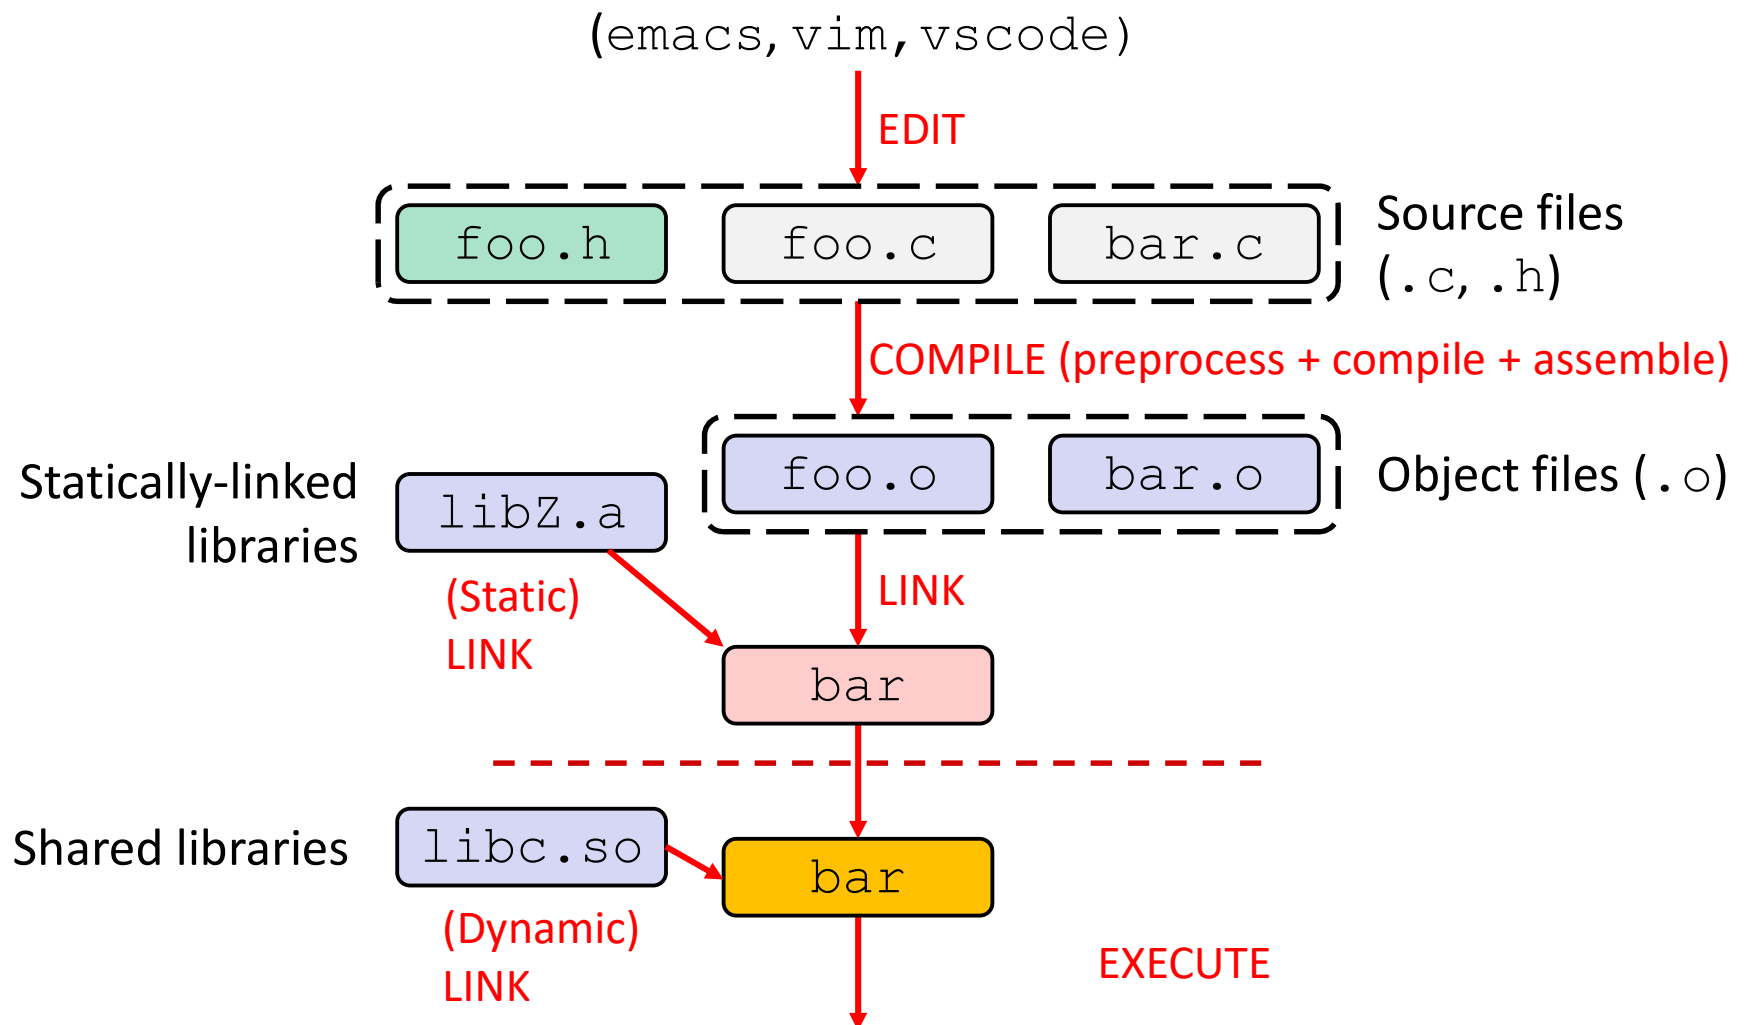
- ❖ What does this mean?

- `argc` contains the number of strings on the command line (the executable name counts as one, plus one for each argument).
- `argv` is an array containing *pointers* to the arguments as strings (more on pointers later)

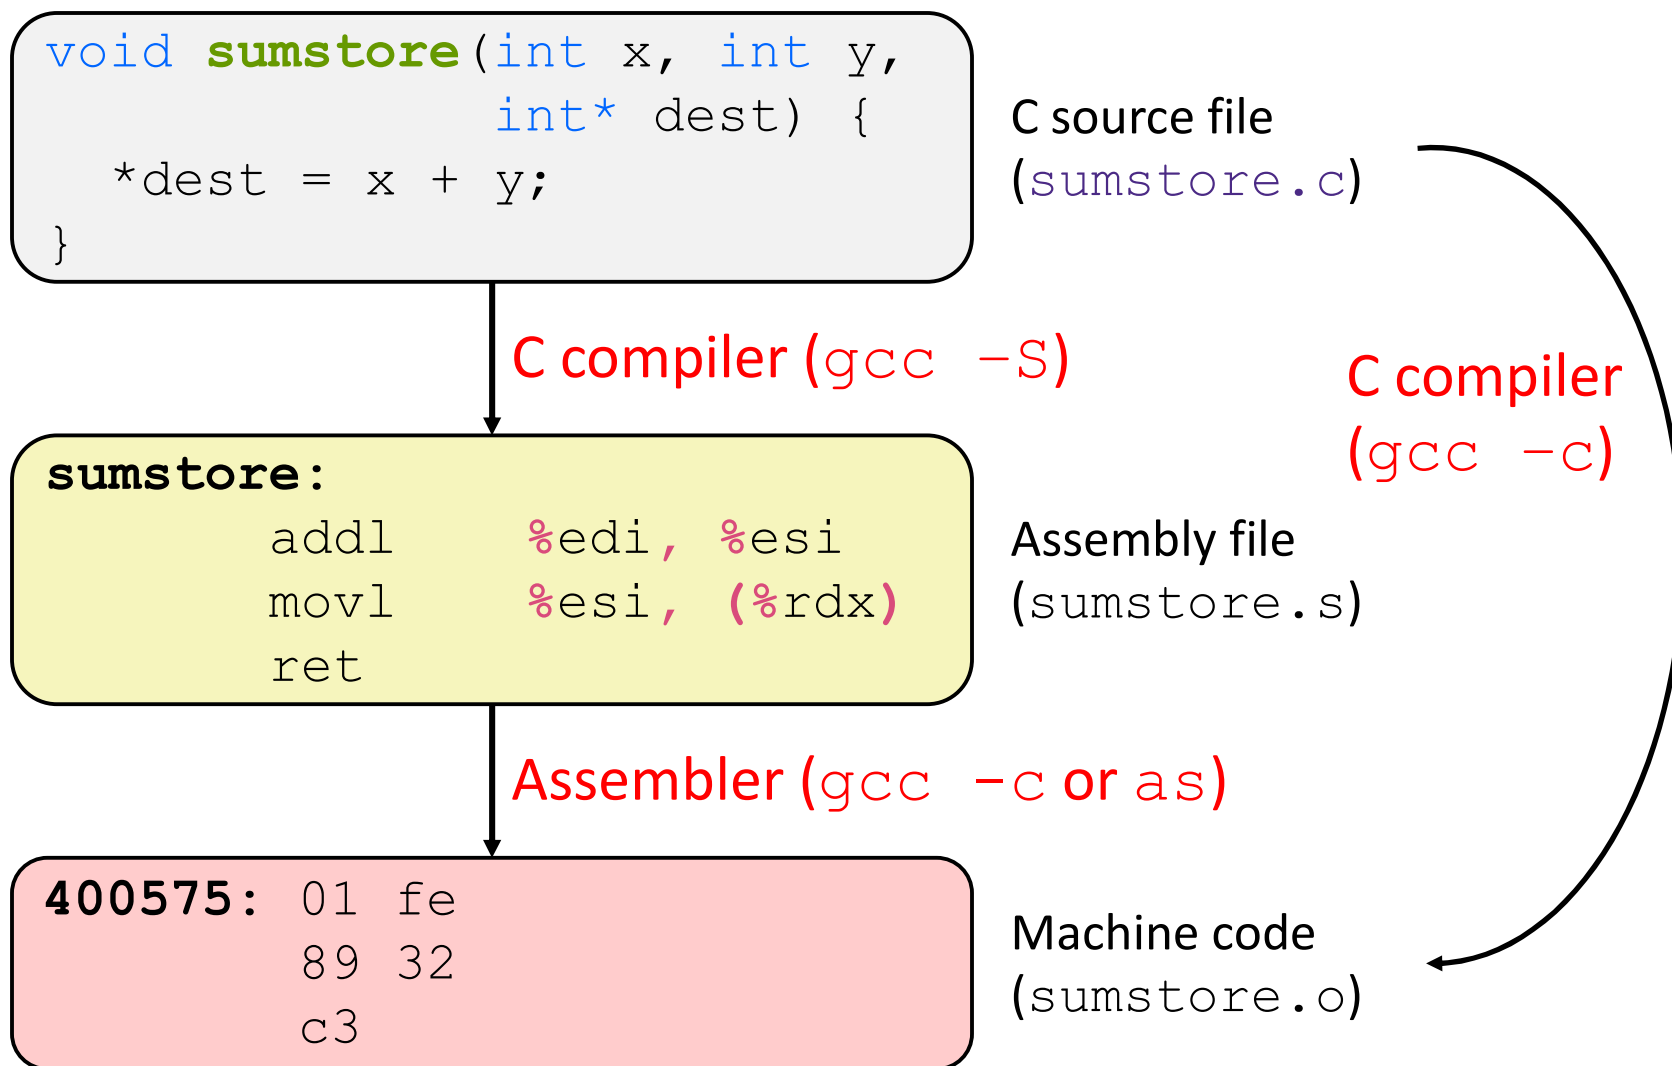
- ❖ Example: `$ foo hello 87`

- `argc = 3`
- `argv[0] = "foo", argv[1] = "hello", argv[2] = "87"`

C Workflow



C to Machine Code



When Things Go South...

❖ Errors and Exceptions

- C does not have exception handling (no `try/catch`)
- Errors are returned as integer error codes from functions
- Because of this, error handling is ugly and inelegant

❖ Crashes

- If you do something bad, you hope to get a “segmentation fault” (believe it or not, this is the “good” option)

Java vs. C (351 refresher)

- ❖ Are Java and C mostly similar (S) or significantly different (D) in the following categories?
 - List any differences you can recall (even if you put 'S')

Language Feature	S/D	Differences in C
Control structures		
Primitive datatypes		
Operators		
Casting		
Arrays		
Memory management		

Java vs. C (351 refresher)

- ❖ Are Java and C mostly similar (S) or significantly different (D) in the following categories?
 - List any differences you can recall (even if you put 'S')

Language Feature	S/D	Differences in C
Control structures	S	
Primitive datatypes	S/D	Similar but sizes can differ (char, esp.), unsigned, no boolean, uninitialized data, ...
Operators	S	Java has >>>, C has ->
Casting	D	Java enforces type safety, C does not
Arrays	D	Not objects, don't know their own length, no bounds checking
Memory management	D	Manual (malloc/free), no garbage collection

Primitive Types in C

❖ Integer types

- `char`, `int`

❖ Floating point

- `float`, `double`

❖ Modifiers

- `short` [int]
- `long` [int, double]
- `signed` [char, int]
- `unsigned` [char, int]

C Data Type	32-bit	64-bit	printf
char	1	1	%c
short int	2	2	%hd
unsigned short int	2	2	%hu
int	4	4	%d / %i
unsigned int	4	4	%u
long int	4	8	%ld
long long int	8	8	%lld
float	4	4	%f
double	8	8	%lf
long double	12	16	%Lf
pointer	4	8	%p

Typical sizes – see `sizeofs.c`

C99 Extended Integer Types

- ❖ Solves the conundrum of “how big is an `long int`?”

```
#include <stdint.h>
```

```
void foo(void) {
```

```
    int8_t a; // exactly 8 bits, signed
```

```
    int16_t b; // exactly 16 bits, signed
```

```
    int32_t c; // exactly 32 bits, signed
```

```
    int64_t d; // exactly 64 bits, signed
```

```
    uint8_t w; // exactly 8 bits, unsigned
```

```
    ...
```

```
}
```

Why do we care how big an int is?

Use extended types in cse333 code

```
void sumstore(int x, int y, int* dest) {
```

```
void sumstore(int32_t x, int32_t y, int32_t* dest) {
```

Basic Data Structures

- ❖ C does not support objects
 - C programs can follow a somewhat object oriented structure, though
- ❖ **Arrays** are contiguous chunks of memory
 - C has a complicated relationship with arrays
 - Arrays have no methods and do not know their own length
 - Can easily run off ends of arrays in C – **security bugs!!!**
- ❖ **Strings** are null-terminated char arrays
 - Strings have no methods, but **string.h** has helpful utilities

```
char* x = "hello\n";
```

x

h	e	l	l	o	\n	\0
---	---	---	---	---	----	----

- ❖ **Structs** are the most object-like feature, but are just collections of fields – no “methods” or functions
 - Support assignment

Function Definitions

❖ Generic format:

```
returnType fname(type param1, ..., type paramN) {  
    // statements  
}
```

```
// sum of integers from 1 to max  
int sumTo(int max) {  
    int i, sum = 0;  
  
    for (i = 1; i <= max; i++) {  
        sum += i;  
    }  
  
    return sum;  
}
```

Why is this a terrible implementation?

Function Ordering

- ❖ You *shouldn't* call a function that hasn't been declared yet

Why?

sum_badorder.c

```
#include <stdio.h>

int main(int argc, char** argv) {
    printf("sumTo(5) is: %d\n", sumTo(5));
    return 0;
}

// sum of integers from 1 to max
int sumTo(int max) {
    int i, sum = 0;

    for (i = 1; i <= max; i++) {
        sum += i;
    }
    return sum;
}
```

Solution 1: Reverse Ordering

- ❖ Frequently used solution; however, imposes ordering restriction on writing functions (who-calls-what?)
 - What if subA calls subB and subB calls subA?

sum_betterorder.c

```
#include <stdio.h>

// sum of integers from 1 to max
int sumTo(int max) {
    int i, sum = 0;

    for (i = 1; i <= max; i++) {
        sum += i;
    }
    return sum;
}

int main(int argc, char** argv) {
    printf("sumTo(5) is: %d\n", sumTo(5));
    return 0;
}
```


Solution 2: Function Declaration

- ❖ Teaches the compiler arguments and return types; function definitions can then be in a logical order

sum_declared.c

Hint: code examples from slides are on the course web for you to experiment with

```
#include <stdio.h>

int sumTo(int); // func prototype

int main(int argc, char** argv) {
    printf("sumTo(5) is: %d\n", sumTo(5));
    return 0;
}

// sum of integers from 1 to max
int sumTo(int max) {
    int i, sum = 0;
    for (i = 1; i <= max; i++) {
        sum += i;
    }
    return sum;
}
```

(Function) Declaration vs. Definition

- ❖ C/C++ make a careful distinction between these two
- ❖ **Definition:** the thing itself
 - *e.g.* code for function, variable definition that “creates” storage
 - (Mostly) must be **exactly one** definition of each thing (no duplicates)
- ❖ **Declaration:** description of a thing
 - It’s the programmer giving the compiler just the information it needs to compile code, but not enough to create the thing being declared
 - *e.g.* function prototype, external variable declaration
 - **Often in header files and incorporated via `#include`**
 - “Must” also `#include` declaration in the code file with the actual definition to check for consistency
 - Needs to appear in **all files** that use that thing
 - Should appear before first use

Multi-file C Programs

definition

C source file 1
(sumstore.c)

```
void sumstore(int x, int y, int* dest) {  
    *dest = x + y;  
}
```

C source file 2
(sumnum.c)

```
#include <stdio.h>  
  
void sumstore(int x, int y, int* dest);  
  
int main(int argc, char** argv) {  
    int z, x = 351, y = 333;  
    sumstore(x, y, &z);  
    printf("%d + %d = %d\n", x, y, z);  
    return 0;  
}
```

declaration

Why is this a terrible way to do this?

Multi-file C Programs Revised

C decl file
(sumstore.h)

```
#ifndef sumstore_h
void sumstore(int x, int y, int* dest;
#endif // sumstore.h
```

C source file 1
(sumstore.c)

```
#include "sumstore.h"
void sumstore(int x, int y, int* dest) {
    *dest = x + y;
}
```




C source file 2
(sumnum.c)

```
#include <stdio.h>
#include "sumstore.h"

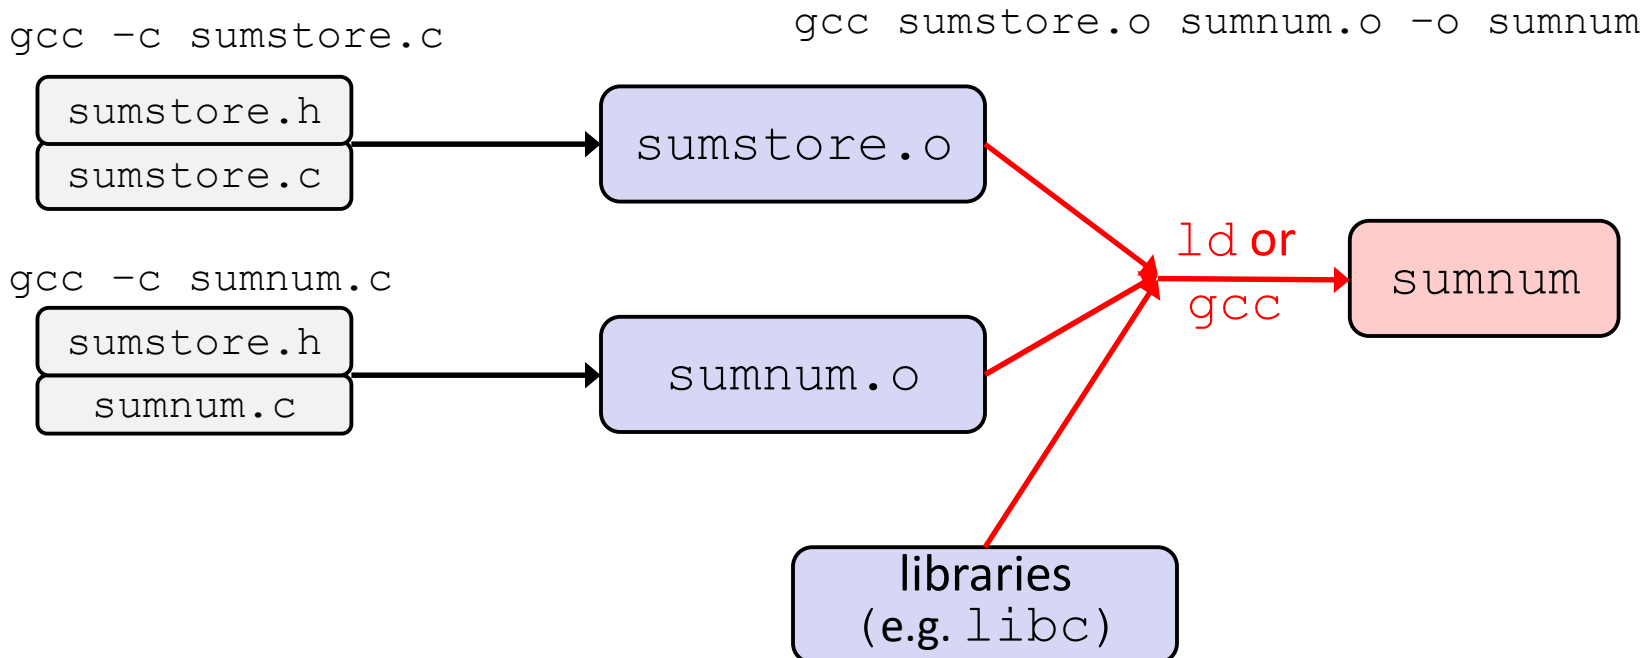
void sumstore(int x, int y, int* dest);

int main(int argc, char** argv) {
    int z, x = 351, y = 333;
    sumstore(x, y, &z);
    printf("%d + %d = %d\n", x, y, z);
    return 0;
}
```



Compiling Multi-file Programs

- ❖ The **linker** combines multiple object files plus statically-linked libraries to produce an executable
 - Includes many standard libraries (*e.g.* `libc`, `crt1`)
 - A *library* is just a pre-assembled collection of `.o` files



To-do List

- ❖ Explore the website *thoroughly*: <http://cs.uw.edu/333>
- ❖ Computer setup: CSE remote lab, attu, or CSE Linux VM
- ❖ Exercise 0 is due 10 am Friday before class
 - Find exercise spec on website, submit via Gradescope
 - Sample solution will be posted Friday after class
 - Give it your best shot to get it done on time
- ❖ Gradescope accounts created just before class
 - Userid is your uw.edu email address
 - Exercise submission: find CSE 333 20au, click on the exercise, drag-n-drop file(s)! That's it!!
- ❖ Project repos created and hw0 out by tonight!!
 - All will become clear in sections tomorrow! 😊